

# Deep Learning

**Recurrent Neural Network**

# Examples of sequence data

Speech recognition



"The quick brown fox jumped  
over the lazy dog."

Music generation

∅



Sentiment classification

"There is nothing to like  
in this movie."



DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AGCCCCTGTGAGGAACTAG

Machine translation

Voulez-vous chanter avec  
moi?



Do you want to sing with  
me?

Video activity recognition



Running

Name entity recognition

Yesterday, Harry Potter  
met Hermione Granger.



Yesterday, **Harry Potter**  
met **Hermione Granger**.

## Motivating example

x: Harry Potter and Hermione Granger invented a new spell.

$$T_x = 9$$

|    |           |           |           |           |       |   |     |           |
|----|-----------|-----------|-----------|-----------|-------|---|-----|-----------|
| y: | 1         | 1         | 0         | 1         | 1     | 0 | 0 0 | 0         |
|    | $y^{<1>}$ | $y^{<2>}$ | $y^{<3>}$ | $y^{<4>}$ | ..... |   |     | $y^{<9>}$ |

 $T_y=9$ 
$$\begin{array}{ll} \mathbf{x}^{(i)<t>} & \mathbf{T}_x^{(i)} \\ \mathbf{y}^{(i)<t>} & \mathbf{T}_y^{(i)} \end{array}$$

# Notation


## Representing words

|  |  |  |  |  |     |  |
|--|--|--|--|--|-----|--|
| x:   |  | Harry Potter and Hermione Granger invented a new spell.  |  |  |     |  |
|  |  | $x^{<1>}$  | $x^{<2>}$  | $x^{<3>}$  | ... | $x^{<9>}$  |
| Vocabulary   |  |  |  |  |     |  |
| $\begin{bmatrix} \text{word1} \\ \text{and} \\ \cdot \\ \cdot \\ \cdot \\ \text{harry} \\ \cdot \\ \text{potter} \\ \cdot \\ \text{wordn} \end{bmatrix}$ | $\begin{bmatrix} 1 \\ 2 \\ \cdot \\ \cdot \\ \cdot \\ 4075 \\ \cdot \\ 6015 \\ \cdot \\ n \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 1 \\ \cdot \\ \cdot \\ \cdot \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ 1 \\ \cdot \\ 0 \end{bmatrix}$ | $\begin{bmatrix} 0 \\ 1 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ 0 \end{bmatrix}$ |     | $\begin{bmatrix} 0 \\ 0 \\ \cdot \\ \cdot \\ \cdot \\ 1 \\ \cdot \\ \cdot \\ \cdot \\ 0 \end{bmatrix}$ |

# RNN

## Why not a standard network?


EACH  $x^{<i>}$   
IS 10000  
dim vector  
because on  
one hot  
vector  
encoding



### Problems:

- Inputs, outputs can be different lengths in different examples.
- Doesn't share features learned across different positions of text.

# Forward Propagation




Bidirectional RNN

He said, "Teddy Roosevelt was a great President."

He said, "Teddy bears are on sale!"

# Forward Propagation




## Simplified RNN notation

$$a^{<t>} = g(W_{aa}a^{<t>} + W_{ax}x^{<t>} + b_a)$$

$$\hat{y}^{<t>} = g(W_{ya}a^{<t>} + b_y)$$

# RNN Cell

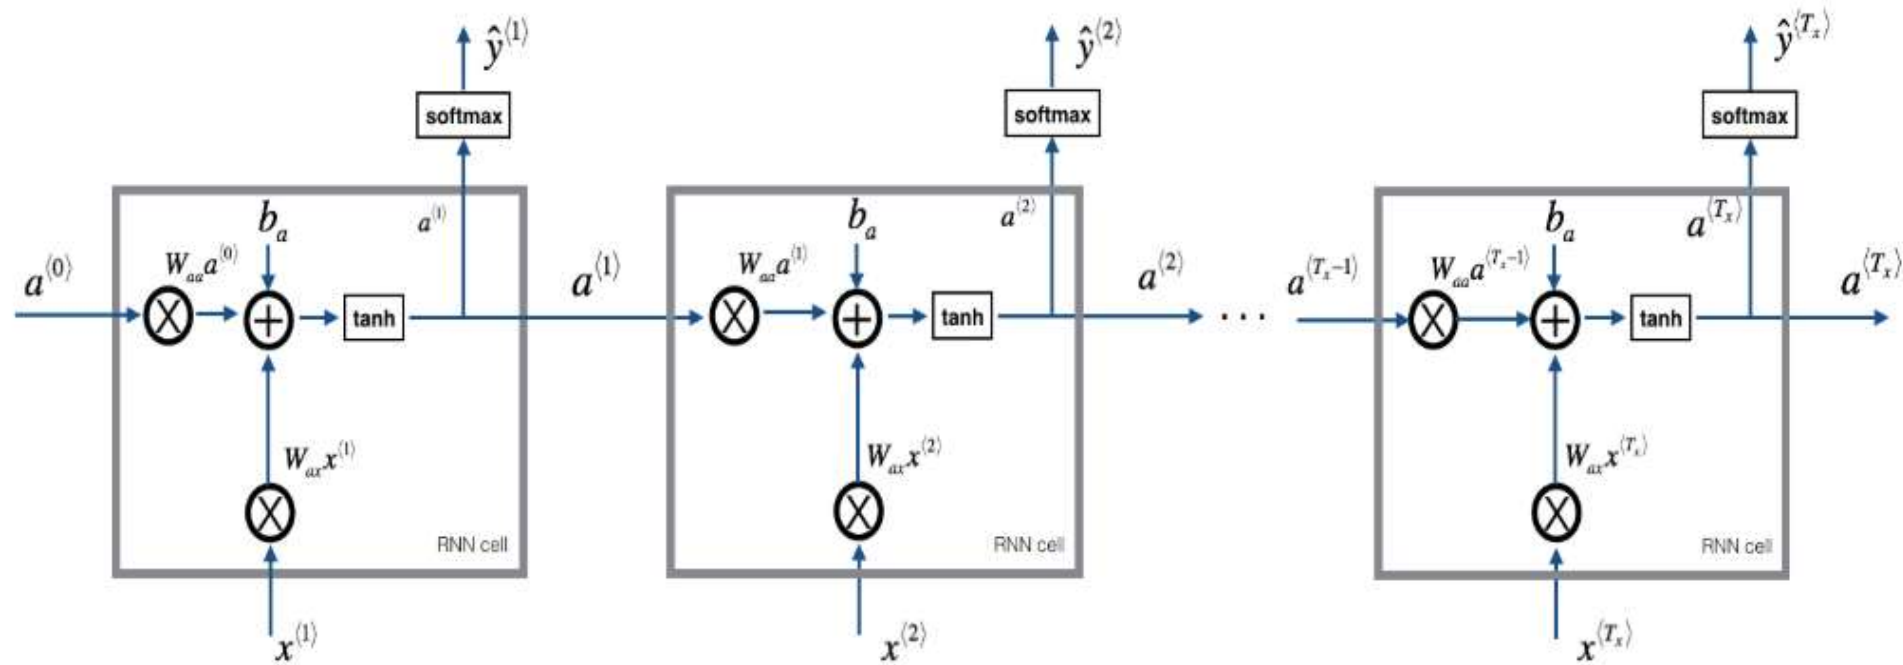


$$a^{(t)} = \tanh(W_{ax}x^{(t)} + W_{aa}a^{(t-1)} + b_a)$$

$$\hat{y}^{(t)} = \text{softmax}(W_{ya}a^{(t)} + b_y)$$



# RNN Forward pass



Basic RNN. The input sequence  $x = (x^{(1)}, x^{(2)}, \dots, x^{(T_x)})$  is carried over  $T_x$  time steps. The network outputs  $y = (y^{(1)}, y^{(2)}, \dots, y^{(T_x)})$

# Example of RNN API

```
rnn = RNN()  
y = rnn.step(x)
```

```
class RNN:
```

```
    # ...
```

```
    def step(self, x):
```

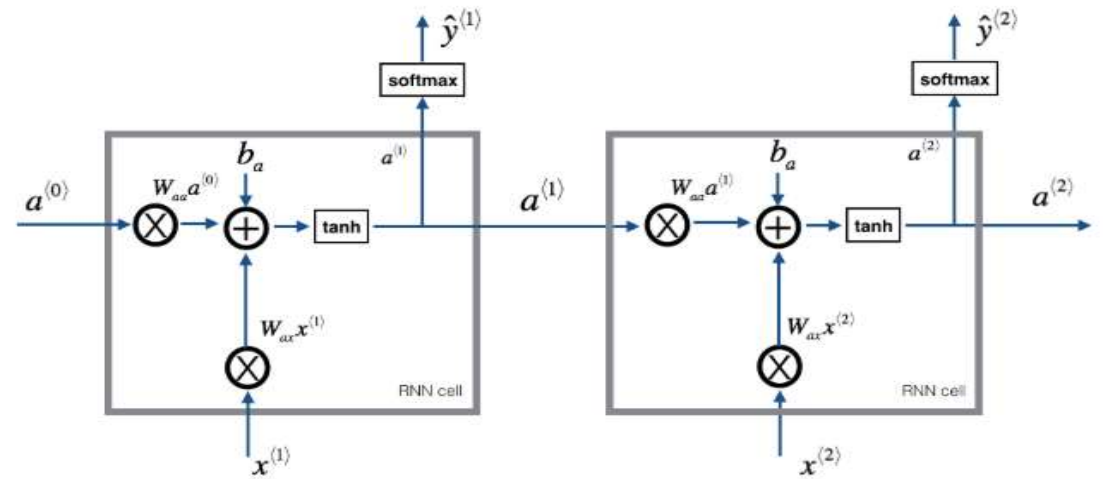
```
        # update the hidden state
```

```
        self.h = np.tanh(np.dot(self.W_hh, self.h) + np.dot(self.W_xh, x))
```

```
        # compute the output vector
```

```
        y = np.dot(self.W_hy, self.h)
```

```
        return y
```



# Example of RNN API

```
class RNN:
```

```
# ...
```

```
def step(self, x):
```

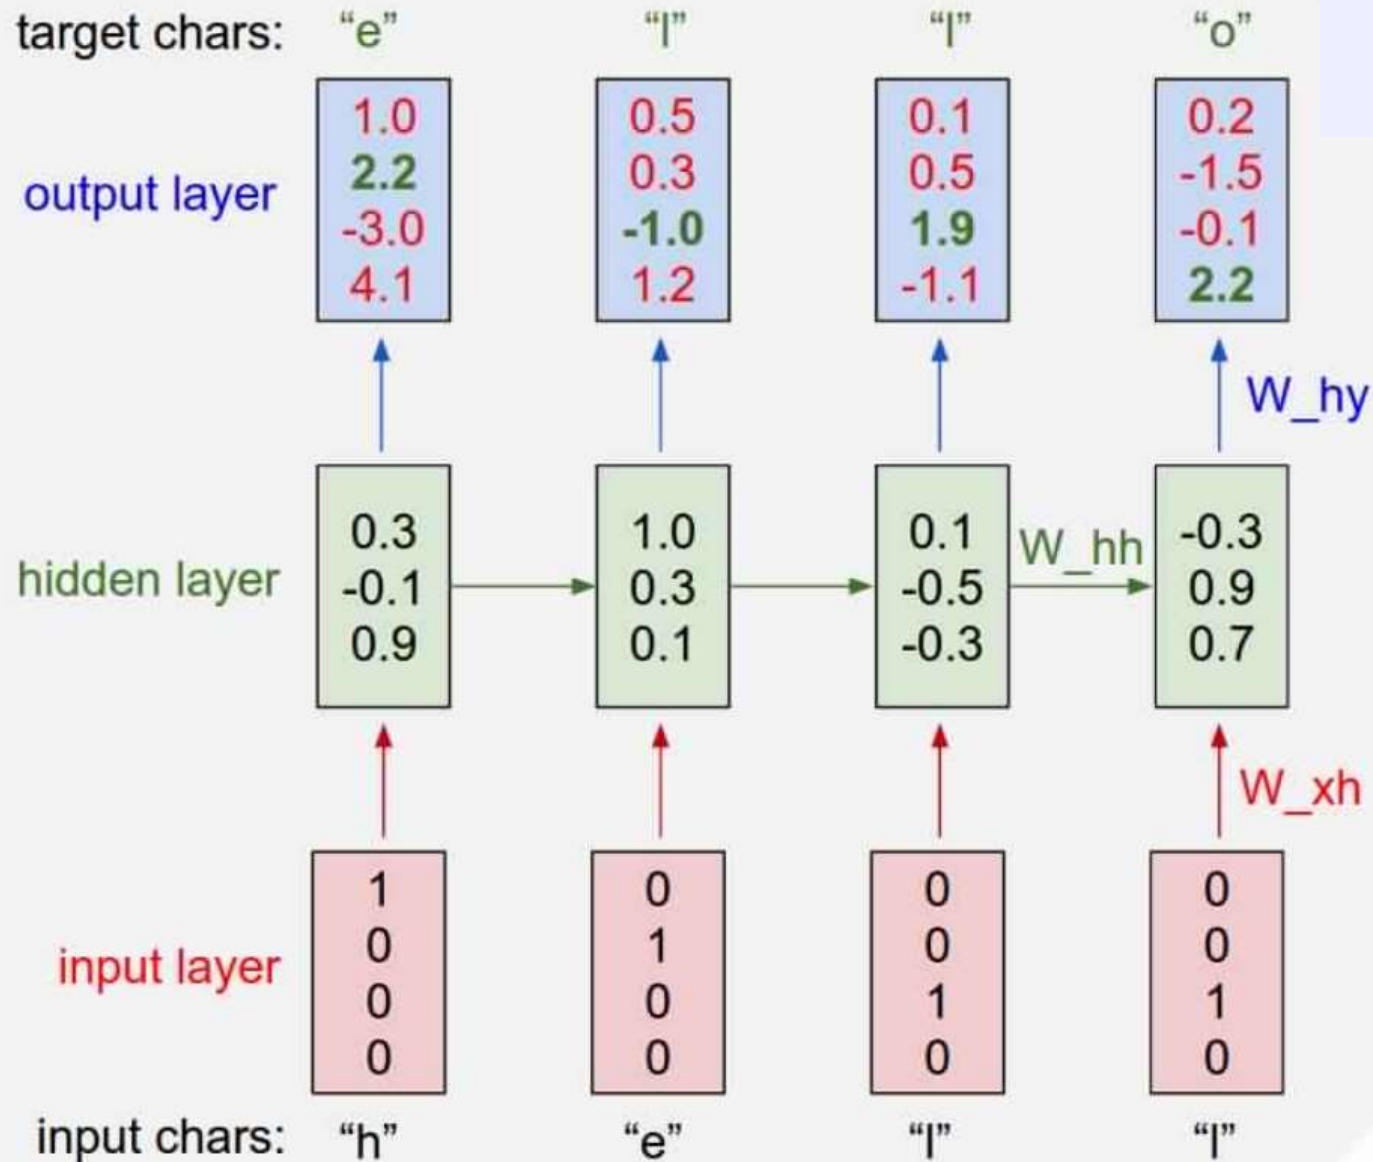
```
    # update the hidden state
```

```
    self.h = np.tanh(np.dot(self.W_hh, self.h) + np.dot(self.W_xh, x))
```

```
    # compute the output vector
```

```
    y = np.dot(self.W_hy, self.h)
```

```
    return y
```



## Simplified RNN notation


$$a^{<t>} = g(W_{aa}a^{<t>} + W_{ax}x^{<t>} + b_a)$$

$$\hat{y}^{<t>} = g(W_{ya}a^{<t>} + b_y)$$


$$g(W_a[a^{<t-1>}, x^{<t>}] + b_a)$$
$$W_a = [W_{aa} \mid W_{ax}]$$

# Backpropagation through time

Forward propagation and backpropagation




# Backpropagation through time




# Different types of RNN


one to one




one to many




many to one



many to many



many to many



# Language model and sequence generation

## What is language modelling?

Speech recognition

The apple and pair salad.

The apple and pear salad.

$$P(\text{The apple and pair salad}) = 3.2 \times 10^{-13}$$

$$P(\text{The apple and pear salad}) = 5.7 \times 10^{-10}$$

$$P(\text{Sentence}) = P(\hat{y}^{<1>}, \hat{y}^{<2>}, \dots, \hat{y}^{<t>})$$



# Language modelling with an RNN

Training set: large corpus of english text.

Tokenization


Cats average 15 hours of sleep a day. <EOS>

$y^{<1>}$     $y^{<2>}$     $y^{<3>}$     $y^{<4>}$    ....    $y^{<8>}$     $y^{<9>}$

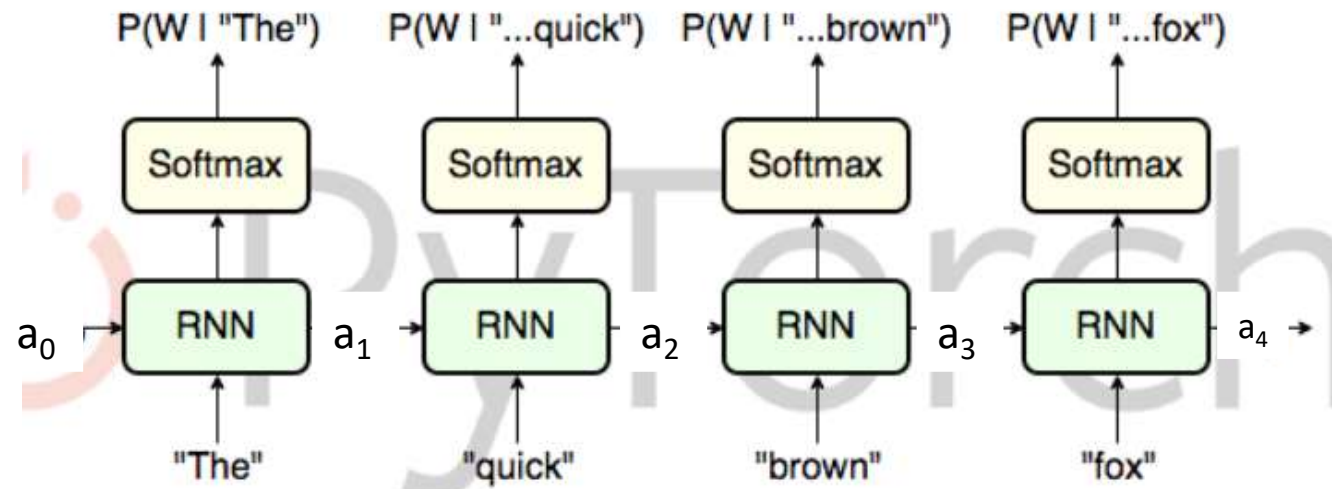
$$x^{<t>} = y^{<t-1>}$$

The Egyptian Mau is a breed of cat. <EOS>

# RNN Language Model




$$\mathcal{L}(\bar{\mathbf{y}}^{<t>}, \mathbf{y}^{<t>}) = - \sum_i y_i^{<t>} \log \bar{y}_i^{<t>}$$
$$\mathcal{L} = \sum_t \mathcal{L}^{<t>}(\bar{\mathbf{y}}^{<t>}, \mathbf{y}^{<t>})$$




$$P(A,B,C,D) = P(A)P(B|A)P(C|A,B)P(D|A,B,C)$$

# Sampling novel sequences

Sampling a sequence from a trained RNN




# Vanishing gradients with RNNs



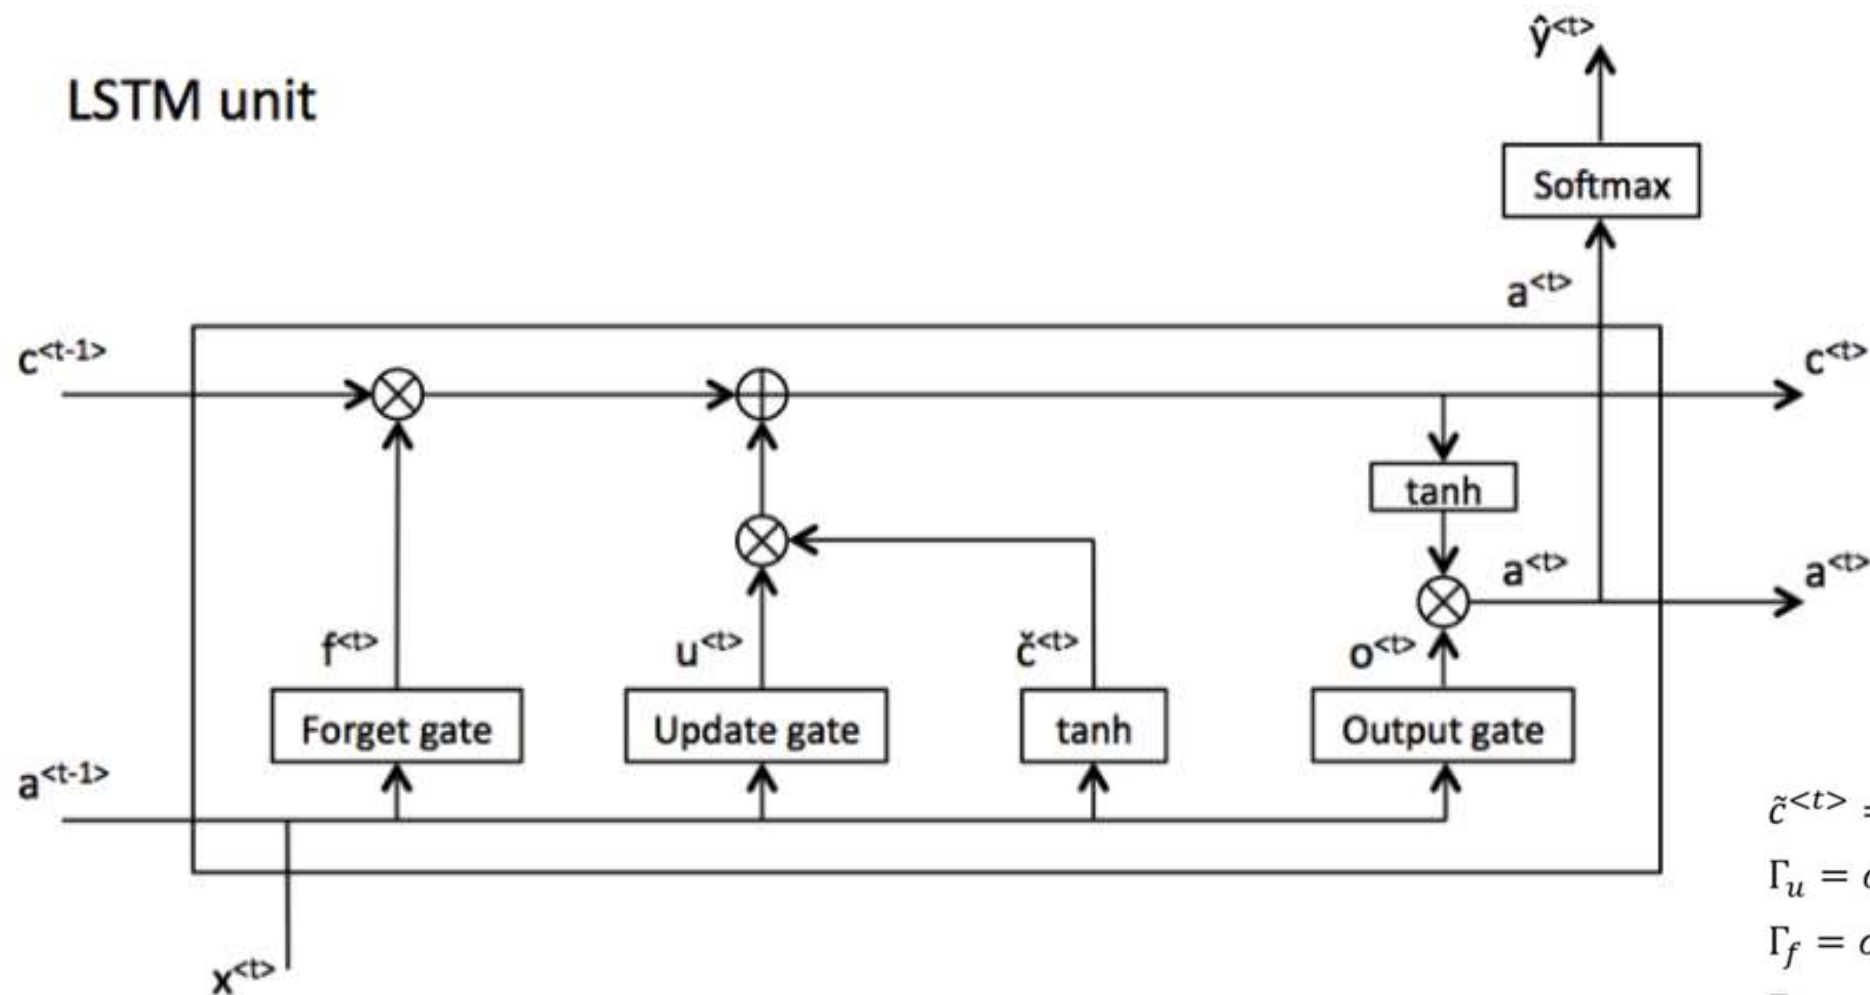
Example 1: The **cat**, which already ate, .....**was** full.

Example 2: The **cats**, which already ate, .....**were** full.

# Local Influence



# LSTM unit



$$\tilde{c}^{<t>} = \tanh(W_c[a^{<t-1>}, x^{<t>}] + b_c)$$

$$\Gamma_u = \sigma(W_u[a^{<t-1>}, x^{<t>}] + b_u)$$

$$\Gamma_f = \sigma(W_f[a^{<t-1>}, x^{<t>}] + b_f)$$

$$\Gamma_o = \sigma(W_o[a^{<t-1>}, x^{<t>}] + b_o)$$

$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + \Gamma_f * c^{<t-1>}$$

$$a^{<t>} = \Gamma_o * \tanh c^{<t>}$$

# LSTM in pictures

$$\tilde{c}^{<t>} = \tanh(W_c[a^{<t-1>}, x^{<t>}] + b_c)$$




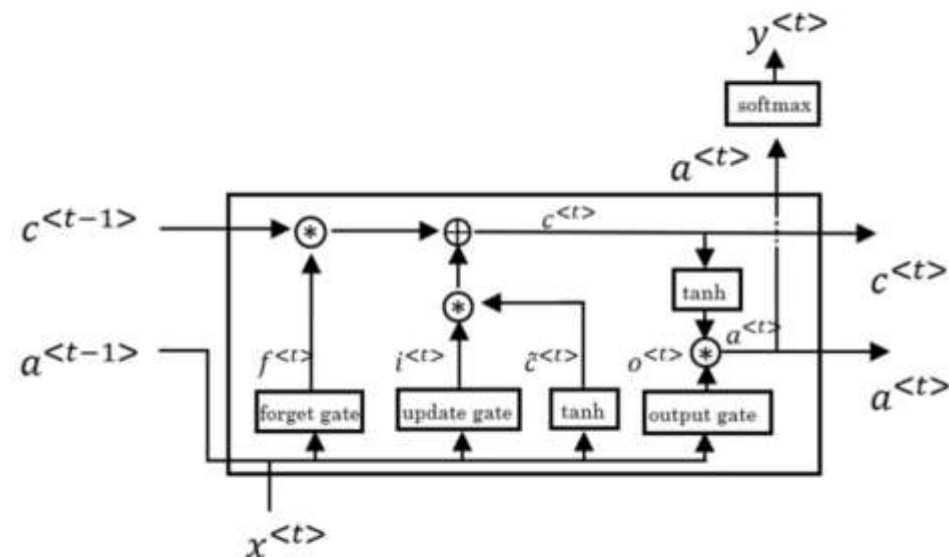
$$\Gamma_u = \sigma(W_u[a^{<t-1>}, x^{<t>}] + b_u)$$

$$\Gamma_f = \sigma(W_f[a^{<t-1>}, x^{<t>}] + b_f)$$


$$\Gamma_o = \sigma(W_o[a^{<t-1>}, x^{<t>}] + b_o)$$

$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + \Gamma_f * c^{<t-1>}$$

$$a^{<t>} = \Gamma_o * \tanh c^{<t>}$$







$$\tilde{c}^{<t>} = \tanh(W_c [c^{<t-1>}, x^{<t>}] + b_c)$$

$$\Gamma_u = \sigma(W_u [c^{<t-1>}, x^{<t>}] + b_u)$$

$$c^{<t>} = \Gamma_u * \tilde{c}^{<t>} + (1 - \Gamma_u) * c^{<t-1>}$$

# Intuition of vanishing gradient in GRU

$$\frac{\partial L}{\partial c^{<t-1>}} = \frac{\partial L}{\partial c^{<t>}} \frac{\partial c^{<t>}}{\partial c^{<t-1>}}$$

and since:

$$\frac{\partial c^{<t>}}{\partial c^{<t-1>}} = 1 - \Gamma_u$$

we have:

$$\frac{\partial L}{\partial c^{<t-1>}} \approx \frac{\partial L}{\partial c^{<t>}} \text{ when } \Gamma_u \approx 0$$


In the notation we are using for gradient descent, this means  $dc^{<t-1>} \approx dc^{<t>}$

# Bidirectional RNN


## Getting information from the future

He said, "Teddy bears are on sale!"

He said, "Teddy Roosevelt was a great President!"

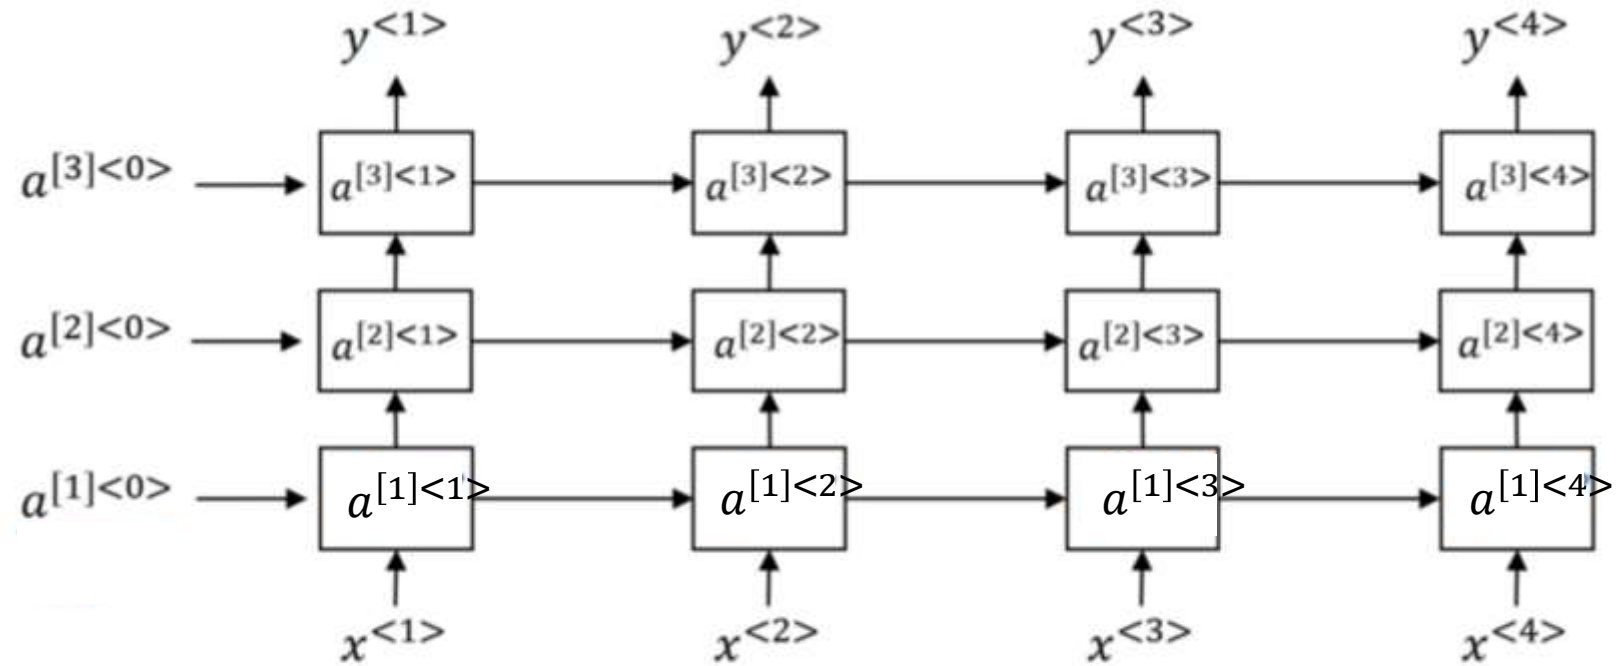


# Bidirectional RNN



<https://towardsdatascience.com/pytorch-basics-how-to-train-your-neural-net-intro-to-rnn-cb6ebc594677>

# Deep RNN



END